

LA ROCHELLE UNIVERSITÉ

LICENSE 3 INFORMATIQUE



Modélisation de bases de données

Project Report: Designing Postgres Database for Genealogy dataset

Professor(s) BOUWMANS Thierry

Members TRAN Minh Duong

PHAM Ngan Ha

NGUYEN Hoang Nhat

Abstract

This project builds a database to organize and study marriage records from France. We take 5,000 marriage documents and create a PostgreSQL database to store them in an organized way. The goal is to design the database so that it is well-organized, easy to use, and can answer questions about families and dates. This report explains how we organized the data and built the database system.

I. Introduction

Old documents and records from France contain important information about families and history. Many of these documents are stored in archives. To study and understand this information, we need to organize it in a database. This project focuses on creating a database system for marriage records from the Vendée region. We show how to take messy information and organize it in a clean, useful way.

I.1. Dataset Overview

We use a file called `mariages_L3_5k.csv` which contains 5,000 marriage records. Each record includes information about two people who got married, their parents, and the place and date of marriage.

The data has these main parts:

- **Record Information:** Each record has an ID number and a date
- **Person A Information:** Name of the first person, and names of their mother and father
- **Person B Information:** Name of the second person, and names of their mother and father
- **Location Information:** The city (commune) and region (department) where the marriage happened

I.2. Why This Project Is Important

When we have many old documents, it is difficult to find information or see patterns. A database makes this easier. Our database has these goals:

1. **Organization:** Store all the information in one place in a clear way
2. **No Duplicates:** Each person should appear only once in the database, even if they appear in many records
3. **Correct Information:** Make sure that all relationships between people and places are correct

II. Database Design & Normalization

II.1. Initial Conceptualization

Initially, we treat the dataset as a single flat table, where the Act ID is the primary key and every other attribute depends on it

Corresponding FD:

```
Act ID → (
    Act Type,
    Person A's Last Name,
    Person A's First Name,
    First Name of A's Father,
    Last Name of A's Mother,
    First Name of A's Mother,
    Person B's Last Name,
    Person B's First Name,
    First Name of B's Father,
    Last Name of B's Mother,
    First Name of B's Mother,
    Town,
    Department,
    Date,
    View Number
)
```

=> Result: 1NF achieved

II.2. Normalization Process

1. For 1NF:

All values are atomic (each cell contains one piece of data).

Each record is uniquely identified by an Act ID => Result: 1NF achieved

2. For 2NF:

Problem: Repetition of data for parents, towns, and departments leads to redundancy.

Solution: Separate entities so that every non-key attribute depends on the whole primary key

Resulting FDs:

- Act_ID → (Type, Date, View_Num, Commune_ID, Person_A_ID, Person_B_ID)
- Person_ID → (Last_Name, First_Name, Father_ID, Mother_ID)
- Commune_ID → (Name, Dept_Code)

=> Result: 2NF achieved

3. For 3NF:

Problem:

- Transitive Dependency: The Department depended on the Town, which depended on the Act.
- Parental names depended on the person

Solution:

- Extracted Department into a standalone table to enforce valid codes (44, 49, 79, 85).
- Established foreign keys (`father_id`, `mother_id`) within the Person table to handle lineage without repeating names.

Final FDs:

`Act_ID` $\rightarrow$ (`Act_Type`, `Date`, `View_Num`, `Commune_ID`, `Person_A_ID`, `Person_B_ID`)

`Person_ID` $\rightarrow$ (`Last_Name`, `First_Name`, `Father_ID`, `Mother_ID`)

`Commune_ID` $\rightarrow$ (`Commune_Name`, `Dept_Code`)

`Dept_Code` $\rightarrow$ (`Valid_Codes`: 44, 49, 79, 85)

=> Result: 3NF achieved

II.3. Final Relation Schema




Figure 1: Database Relation Schema

Cardinality:

- Each parent can have more than one child but a child can only have one of each parent.
- Each department can have more than one town, but a town must only belong to one department.

III. Data Extraction & ETL Process

To get our data from the messy CSV file into our clean database, we created a two-step process using Python scripts. The first script, `utils/data_parser.py`, handles the cleaning and organizing, and the second script, `main.py`, loads it into the database.

III.1. Extraction

First, we pull all the data from the `mariages_L3_5k.csv` file. We used the pandas library in Python to help with this. An important step was to read everything as text first, so we didn't accidentally change data, like dropping the '0' from a department code like '05'. We also made sure to use UTF-8 encoding, which helps read special French characters correctly.

III.2. Transformation

This is where we do most of the cleanup work. Our script goes through the raw data row by row and gets it ready for our normalized tables.

Data Cleaning

We made a simple cleaning function that does two things:

- It removes any extra spaces from the beginning or end of each piece of data.
- It finds any empty or "n/a" values and marks them as properly missing, so we can handle them consistently.

Finding and Removing Duplicates

To avoid storing the same person or town multiple times, our script keeps track of what it has already seen.

- **Departments and Towns:** It checks for valid department codes and skips any rows with bad data. It makes sure a town is unique by looking at both its name and its department code.
- **People:** To decide if a person is unique, we look at their full name and also who their parents are. This helps tell apart two people who might have the same name.

Fixing Formats

We also had to fix a few things to match our database design:

- **Dates:** We converted the date from the "DD/MM/YYYY" text format into a real date that the database can understand.
- **View Number:** We took columns like "48/210" and just kept the first number (48), which is the page number we care about.


III.3. Loading

Instead of manually importing the files, we used a second Python script, `main.py`, to automatically load our cleaned data into the PostgreSQL database. This script uses a library called SQLAlchemy, which helps Python talk to databases.

The loading process is done carefully to avoid errors:

1. **Final Check:** Before loading, the script does one last check on our cleaned data tables to make sure all the ID numbers are valid integers and that no essential information is missing.

2. **Order of Operations:** It loads the tables in the correct order: first **department**, then **town** (which needs departments to exist), then **person**, and finally **act** (which needs towns and people).
3. **Handling People:** The **person** table was a bit tricky because a person's parents are also in the same table. To solve this, we first insert all people with their parent fields empty. Then, once everyone is in the table, we go back and add the parent-child links.
4. **Batching:** To make it faster, the script inserts the data in large chunks (batches) instead of one row at a time.
5. **Updating Sequences:** After inserting everything with specific IDs, we reset the database's automatic ID counters. This ensures that if we add new data later, we won't have any ID conflicts.

A terminal window with a dark background. The title bar shows the file path ".../postgres-genealogy-project" and several window controls. The terminal content shows the command "uv run python main.py" being executed. The output indicates that 5000 rows were loaded from a CSV file, with two rows (1684 and 2478) having unparseable dates. It then shows a summary of the data loaded into the PostgreSQL database: 4 rows for department, 375 rows for town, 12271 rows for person, and 5000 rows for act. The terminal ends with "Loaded data into PostgreSQL."

```
.../postgres-genealogy-project ? main x! ? o base v3.12.7 22:23
> uv run python main.py
Loaded 5000 rows from raw_data/mariages_L3_5k.csv
  △ Row 1684: unparseable date '2?/11/1781'
  △ Row 2478: unparseable date '1628-1629'

Parsed successfully (0 rows skipped):
  department:      4 rows
    town:      375 rows
  person:  12271 rows
    act:    5000 rows
Loaded data into PostgreSQL.
```

Figure 2: Terminal output after running the ETL pipeline with main.py

IV. Genealogical Queries & Results

Here are the SQL queries designed to answer specific genealogical questions based on our database schema, along with their results.

1. The number of communes per department

```
SELECT dept_code, COUNT(*) AS total_town
FROM town
GROUP BY dept_code;
```

2. The number of acts in Luçon

```
SELECT COUNT(*)
FROM act
JOIN town ON act.town_id = town.id
WHERE town.name ~* '^LUÇON$';
```

Query	Query History
45	-- QUERY FOR QUESTIONS --
46	-- 1. La quantité de communes par département
47	-- Number of towns per department
48	EXPLAIN ANALYZE
49	SELECT dept_code, COUNT(*) AS total_town
50	FROM town
51	GROUP BY dept_code;

Data Output	Messages	Notifications
	dept_code character varying (2)	total_town bigint
1	85	313
2	49	2
3	44	9
4	79	51

(a) Result for Query 1

Query	Query History
53	-- 2. La quantité de actes à LUÇON
54	-- Number of acts that belong to the town LUÇON
55	SELECT COUNT(*)
56	FROM act
57	JOIN town ON act.town_id = town.id
58	WHERE town.name ~* '^LUÇON\$';

Data Output	Messages	Notifications
	count bigint	
1	90	

(b) Result for Query 2

3. The number of marriage contracts before 1855

```
SELECT COUNT(*)
FROM act
WHERE act.act_type = 'Contrat de mariage'
AND act.act_date < '1855-01-01';
```

4. The town with the highest number of marriage publications

```
SELECT town.name, COUNT(*) AS total_publications
FROM act
JOIN town ON act.town_id = town.id
WHERE act.act_type = 'Publication de mariage'
GROUP BY town.name
ORDER BY total_publications DESC
LIMIT 1;
```

Query	Query History
59	-- 3. La quantité de "contrats de mariage" avant 1855
60	-- Number of marriage contracts before 1855
61	SELECT COUNT(*)
62	FROM act
63	WHERE act.act_type = 'Contrat de mariage'
64	AND act.act_date < '1855-01-01';
65	
66	
67	-- 4. La commune avec la plus quantité de "publication

Data Output	Messages	Notifications
	count bigint	
1	196	

(a) Result for Query 3


Query	Query History
67	-- 4. La commune avec la plus quantité de "publica
68	-- Town with the most "publications de mariage"
69	SELECT town.name, COUNT(*) AS total_publications
70	FROM act
71	JOIN town ON act.town_id = town.id
72	WHERE act.act_type = 'Publication de mariage'
73	GROUP BY town.name
74	ORDER BY total_publications DESC
75	LIMIT 1;
76	

Data Output	Messages	Notifications
	name character varying (255)	total_publications bigint
1	SAINT PIERRE DU CHEM...	20

(b) Result for Query 4

5. The date of the first and last records

```
SELECT MIN(act_date) as first_act, MAX(act_date) as last_act
FROM act;
```




	first_act date	last_act date
1	1581-12-23	1915-09-14

Figure 5: Result for Query 5

Summary of All Query Results

The following image shows the combined results of all queries executed through the Python script.



```
.../postgres-genealogy-project  main x!  base v3.12.7  22:08
> uv run python query.py
Q1 Number of communes per department: [{'dept_code': '85', 'total_town': 313}, {'dept_code': '49', 'total_town': 2}, {'dept_code': '44', 'total_town': 9}, {'dept_code': '79', 'total_town': 51}]
Q2 Number of acts in Luçon: 90
Q3 Number of marriage contracts before 1855: 196
Q4 Town with the highest number of marriage publications: SAINT PIERRE DU CHEMIN
Q5 Date of the first and last records: (datetime.date(1581, 12, 23), datetime.date(1915, 9, 14))
```

Figure 6: Summary of All Query Results from Python Script

V. Bonus: Scalability & Noisy Data

TODO:

Challenge Overview: Dealing with the 564k register mariages.csv file .

Difficulties Encountered: Description of the "noise" found and how you resolved it.

Performance Comparison: Comparison of query results between the 5k and 564k datasets.

VI. Conclusion